

RECOMMENDER SYSTEM CHALLENGE 2022-23

Alessandro Rossi - 10649607
Davide Zanutto - 10938003



OUR PROCESS

Trying to understand the **dataset** and which algorithms worked best

TRIAL & ERROR

BEGINNING

1° DEADLINE

BASIC HYBRID

One of our first **hybrid** recommenders: just slim elastic net & rp3Beta

Tuning of hybrids and single algorithms to find the best **combinations**

TUNING & TESTS


MAIN PHASE

2° DEADLINE

FINAL HYBRID

Hybrid that worked best composed by **6 different** algorithms

UNDERSTANDING THE DATA



URM

- 41.6k Users
- 24.5k Items
- 233k Impressions

ICM Type

- Types of items
- From 1 to 7
- Not present for every item

ICM Length

- Number of series' episodes
- Not present for every item


BASE MODEL EVALUATIONS

Performances: Valid MAP@10

SLIM-EN Map@10: 0.036425 01	EASE_R Map@10: 0.035487 02	RP3_BETA Map@10: 0.033997 03	P3ALPHA Map@10: 0.031520 04
ITEMKnnCF Map@10: 0.031182 05	IALS Map@10: 0.026376 06	SLIM-BPR Map@10: 0.025328 07	PureSVD Map@10: 0.019294 08

Where to start?

- Convert user and interaction csv into a binary URM
- Train-Validation Split (80:20)
- Understand course's repository
- Start from simpler solutions
- Don't panic



Our specs for the project

- Used Kaggle's integrated environment:
 - 4 Core CPU
 - 30 GB RAM
 - 12 hours max execution time for CPU


HYBRIDATION METHODS

STACKING

ICM type stacked vertically to the URM

PREDICTION COMBINATION

Using a weighted average



MERGING SIMILARITY MATRIX

With models that use the same structure

ROUND ROBIN

Combination of csv prediction files
It didn't go well...

First Deadline Hybrid and Improvements

FIRST DEADLINE HYBRID

Our first best hybrid was a similarity merge between SLIM-EN and RP3 beta

Unfortunately it was not enough to pass the last baseline so we were doing something wrong (or maybe we had been too much optimistic)

So the next step was to understand what were we doing wrongly...

Improvement #1

Stacking ICM



We noticed that some base models improve their performances by stacking the icm

Improvement #2

TUNING



By exploiting Kaggle's potential we managed to run up to 20 different tuning at the same time


Improvement #3

HYBRID MORE




After the tuning phase we managed to pass the last baseline so we focused on improving our hybrid model

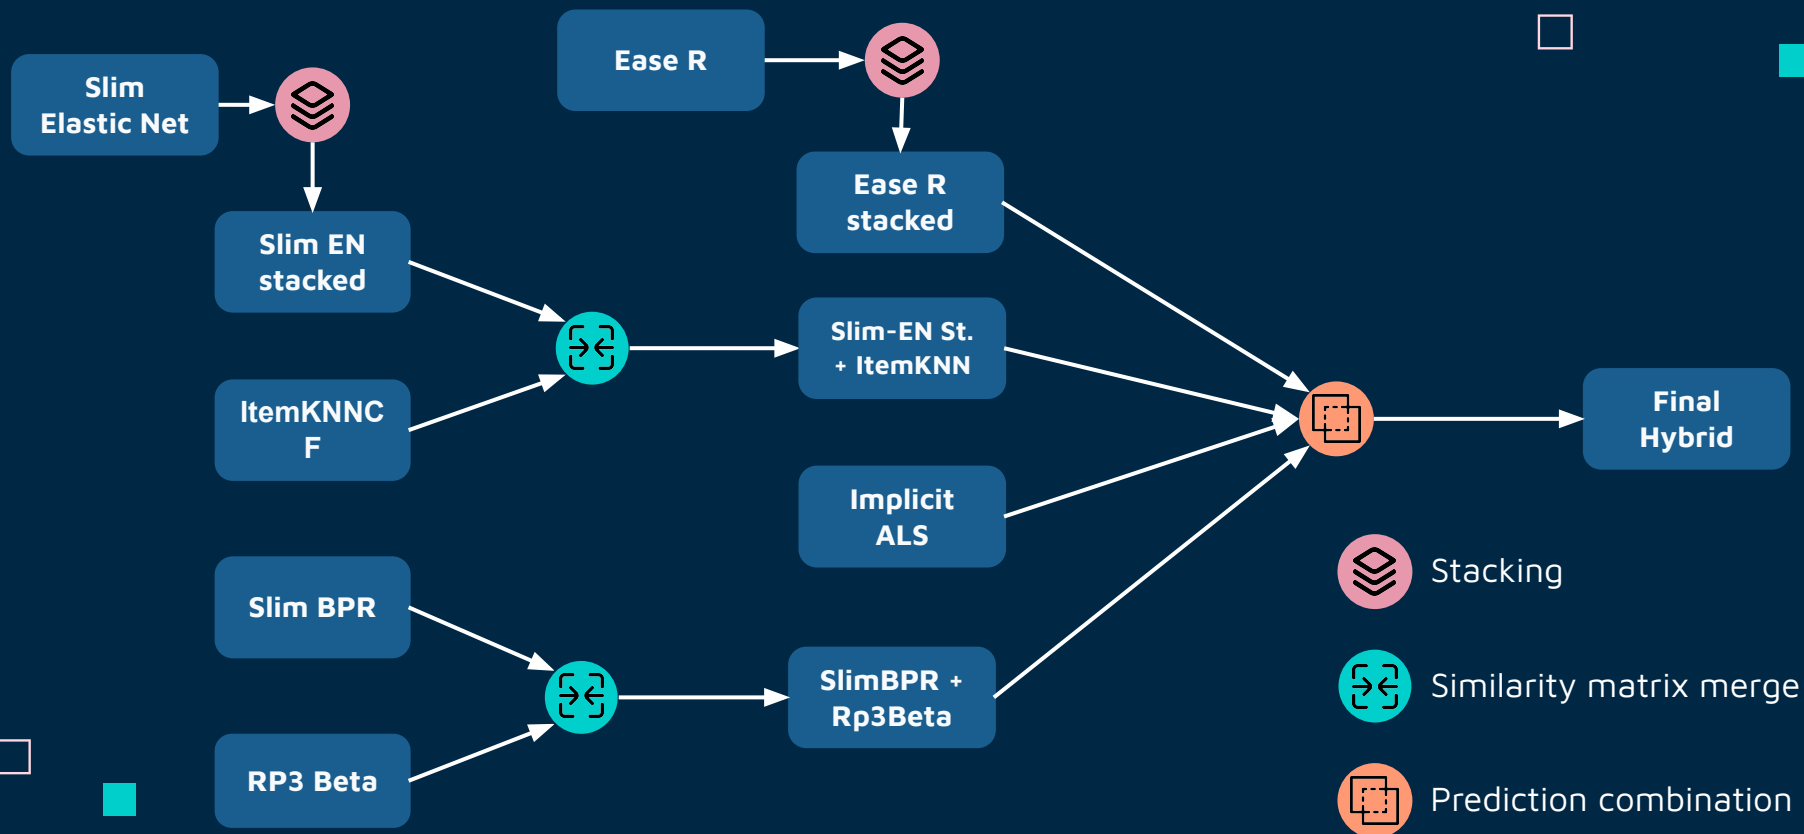
When you run 10 tuning at the same time but none of them improves



HYPERPARAMETER TUNING



OUR BEST HYBRID



WHAT DIDN'T WORK

- Building a **not-binary** URM exploiting information from ICMs
- Some algorithms did good individually but bad when hybridized (Ex: P3Alpha)
- List merging (Round Robin)
- Using Top Popular recommendations for specific users
- Stacking with ICM_length
- Tailor different models for specific portions of the data

OUR FINAL RESULT

- Thanks to the mentioned improvements and the final hybrid, we managed to reach an online map of: 0.06064 (private leaderboard)



OUR TEAM



Davide Zanullo



Data Science student
@ EIT Digital Master School
(POLIMI, TU/e)



Alessandro Rossi

Data Science Student
@ EIT Digital Master School
(POLIMI, UPM)



THANK YOU
FOR YOUR ATTENTION

